

DOCUMENT TYPE: Guidelines

CONTACT: Evan Tachovsky, Chief Data & AI Officer. Lauga Malaguzzi - Global Director, Research Integrity.

OWNER: Research, Data and Innovation / !General

SCOPE: Institutional

EFFECTIVE: 9/18/2023

REVISED: 4/28/2026

STATUS: Approved

APPROVED BY: Research, Data and Innovation

Generative AI Use at WRI: Guidelines for Programmatic Use

1. POLICY

This document is a guide to using Generative AI to support research, communications and programmatic delivery at World Resources Institute (WRI).

Scope

This document offers guidance on overarching principles, current use cases, ways to get support with AI, and answers Frequently Asked Questions. For each use case, we provide the following color-coded framework:

- **Green Use Cases:** Use Generative AI with appropriate care and **always follow three rules: (1) keep a human in the loop to review the outputs of the AI; (2) do not share sensitive, confidential, or private information in your prompts to the AI; and (3) do not outsource judgement or taste to AI.**
- **Yellow Use Cases:** Using Generative AI for this task can be helpful, but only if the user is an expert in the field or topic they are exploring and has the editorial experience to identify and correct gaps and mistakes introduced by AI.
- **Red Use Cases:** Using Generative AI for this task may cause more harm than good. We advise against using Generative AI for red use cases currently.

The guidance also includes role-specific **best practices for researchers, communications professionals, and data, product, and engineering staff.**

For questions or comments, please email Evan Tachovsky (Chief Data & AI Officer) and Laura Malaguzzi (Global Director, Research Integrity).

2. CONTEXT

INTRODUCTION

This document is a guide to using Generative AI to support research, communications and programmatic delivery at World Resources Institute (WRI). Generative AI refers to algorithmic systems designed to produce text, images, or code that is indistinguishable from content created by humans. The guide primarily focuses on large language models (LLMs) and AI platforms such as Gemini, ChatGPT, Claude, Llama, Mistral, or DeepSeek as well as image generators such as DALL-E or Midjourney. This guide does not apply

to traditional machine learning and AI models such as those used routinely at WRI for classification or prediction tasks, e.g. to classify satellite imagery or predict a risk score.

The document is divided into four sections:

- A short introduction to Generative AI as a technology.
- A set of principles and practices for responsible use of Generative AI at WRI
- Guidance for specific use cases relevant to WRI's work.
- How to get help with Generative AI at WRI.

HOW DOES GENERATIVE AI WORK?

Generative AI tools are complex systems composed of *data*, machine learning models, and platforms used by people. At heart, Generative AI tools are fundamentally prediction machines, and the processes under the hood are complex, variable, and probabilistic.

Generative AI starts with massive amounts of *data*. The current generation of tools are trained on essentially all the data on the public internet and in libraries plus a bevy of private data sources.

Generative AI *models* use machine learning algorithms to generate content by predicting the next token, or combination of tokens, in a sequence based on patterns learned from their training data. These models are trained in two phases: first, an intensive pre-training phase where the model absorbs general patterns and knowledge from the training data and arranges and represents it in billions of neural network parameters. These are probability-based maps of associations. These multidimensional neural networks reflect links between words and distances between them. In the second phase, models are refined using methods like supervised fine-tuning and reinforcement learning from human feedback to tailor the model's performance¹.

More recently, a third category of models called "reasoning models" has emerged. These are designed to think through answers step-by-step before responding, which generally improves performance on analytical tasks. Once a model is trained, it uses a process called inference to produce responses by predicting and generating tokens, which are then presented to the user as chunks of text or code. Image generators generally use a parallel but different process called diffusion to produce images or video.

Generative AI *platforms* are the websites and applications that most people use to access the models. These platforms can be generic, like ChatGPT, or focused on a specific task like Cursor for coding or Elicit for scientific literature reviews. Increasingly, platforms also offer "agentic modes", which allow AI models to take a sequence of actions, for example browsing the web or writing and deploying code, with minimal human intervention. Note that many engineers use AI models and agents directly from their command line. They can install tools like Claude Code and execute commands by typing them directly into a terminal, skipping the graphic user interface. In contrast, rather than interacting directly with the underlying data or

¹ For a visual explain of model architecture for generative AI, see this helpful explainer from the Financial Times: <https://ig.ft.com/generative-ai/>

models general users almost always use platforms featuring windows, icons and buttons². Beyond standalone AI platforms, Generative AI is increasingly embedded directly into existing business tools as a feature, for example AI-integrations in Salesforce or Asana. These embedded AI tend to support specific tasks like summarization or automated notifications.

WHAT DO WE KNOW ABOUT GENERATIVE AI'S EFFICACY?

The combined functionality of Generative AI systems is still best described as a “**jagged frontier**” (Dell’Acqua et al. 2026). The systems excel in certain tasks while underperforming in others, often in unintuitive ways. While the literature on AI effectiveness and risk is evolving rapidly, a consensus supports the following findings:

- **AI generally improves task-specific productivity, particularly in writing, coding, and customer service.** Noy and Zhang (2023) report substantial improvements on writing tasks, with the greatest benefit for lower-skilled writers. Cui et al. (2026) find that LLMs speed software development with the highest rates of use and productivity gains among lower skilled developers. And Brynjolfsson et al. (2025) observed productivity gains among customer support agents using AI tools with new workers benefiting the most.
- **AI systems are resource intensive.** They currently use more resources—i.e. electricity and water— than other types of searches or computational activity (see e.g. Saenko 2023; Dodge et al. 2023). This resource use intensity is concentrated in specific geographies and still not fully understood because of gaps in available data (Chen 2025). Several researchers are [benchmarking electricity use of the different models](#).
- **Challenges with hallucinations and sycophancy persist.** Ji et al. (2023) and Azaria et al. (2023) highlight how AI hallucinations frequently arise from limitations in training data, insufficient grounding mechanisms, overconfidence in probabilistic outputs or combining unlike or incommensurable data. Malmqvist (2025) found that sycophantic behavior, where models prioritize user approval over accuracy, was likely the result of reinforcement learning from human feedback, which rewarded agreeability over truthfulness. Despite meaningful improvement, both sycophancy and hallucinations around important content like citations remain problematic.
- **Tradeoffs between automation and verification.** Wang et al. (2025) emphasize the importance of verifying output for each task. When AI agents work autonomously across many tasks, the necessary verification time may exceed any initial time savings.

There is emerging consensus around the following findings:

- **Deskilling is a risk for AI users.** There is emerging evidence to suggest that as more tasks are delegated to automated systems, human knowledge, critical thinking, and specialized skills

² Note that the user-platform names are often used interchangeably with underlying model names. For reference as of April 2025: Google unified its platform and model under the Gemini name and is currently on the Gemini 3.x model series; OpenAI’s ChatGPT platform is powered by the GPT-5 series models while Microsoft Copilot also runs OpenAI models; Anthropic’s Claude platform is powered by the Claude-4 model family including Sonnet and Opus models; Meta’s Llama-4 is a standalone open weight LLM; Mistral AI offers a family of open-weight models accessible via its platform; and DeepSeek is platform that leverages a range of models including DeepSeek R1 and V3.

could erode (Acemoglu et al. 2026). Research in medical practice has shown evidence of deskilling for doctors who overly rely on AI diagnostic tools (Budzyń et al. 2025) and there are similar findings with software developers who intensely rely on AI assistants (Shen and Tamkin, 2026).

- **AI systems are vulnerable to adversarial attacks and bad faith uses.** Zou et al. (2023) demonstrate that large language models are susceptible to prompt injection and adversarial attacks, posing reliability challenges.
- **AI is an ethical and governance challenge.** Whittlestone and Clarke (2023) discuss the complexities in governing AI applications within academic and research settings, emphasizing the need for robust frameworks to address ethical and operational concerns. Ethical challenges for WRI are included in Table 1 and span plagiarism, privacy concerns, bias and use of resources.
- **AI poses information integrity risks for society at large and knowledge work specifically.** AI significantly lowers the barriers for generating and distributing content at scale. This in turn blurs the line between fact, opinion, and fiction, increasing the potential for large-scale mis- and disinformation campaigns (NIST, 2024). There are also extensive concerns about non-transparent integration of third-party data and assets across AI supply chain, which make it hard for downstream users to trace provenance or quality of AI outputs.

AI OPPORTUNITIES AND RISKS FOR WRI

This new technology presents novel opportunities and challenges for staff across WRI – see Table 1 for specific risks and opportunities for WRI staff.

Table 1. Common risks and opportunities of using Generative AI at WRI

Risks	Opportunities
There are risks that generative AI will: • Provide inaccurate information that <i>appears</i> superficially correct. • Contribute to deskilling, reducing experts' abilities over time. • Use excessive energy and water, contributing to climate change and environmental degradation. • Introduce or perpetuate bias from underlying training data. • Expose sensitive information to others. • Plagiarize and/or violate copyrights by directly copying underlying data • “Disintermediate” by reducing direct access to our work and reducing our ability to choose how our work is presented and to whom. • “Disrupt” by shifting the location of work to AI models and tools and away from organizations like ours.	There are opportunities for generative AI to: • Expand the capacity of small teams, by serving the role of an eager research assistant for some tasks. • Speed up time-consuming, repetitive tasks. • Deepen insight and understanding of complex literature. • Generate new ideas or content that can be further explored and developed. • Audit your quantitative research for replicability. • Develop an opposing view to help strengthen your arguments. • Customize and optimize work to appeal to and reach particular audiences. • Enable us to connect more directly or even integrate into research and rulemaking processes in our audiences.

PRINCIPLES TO USE GENERATIVE AI RESPONSIBLY

We operate using a set of [Responsible AI Principles and Practices](#) that are grounded in our [Institutional Values](#) and have been refined through over a decade of institutional innovation and data science work:

Responsible AI Principles

1. **Be curious:** We encourage our staff to base decisions on evidence – to design experiments and measure results. Similarly, when it comes to new technologies we should neither trust the hype around them nor avoid them altogether, but move forward only when the data shows potential for positive impact and manageable risk. Most experiments fail; some succeed and progress; all produce valuable learning.
2. **Be user centered:** We build new technology to solve problems for people, nature and the climate. We do not build technology for novelty's sake or in ways that create large risks for those we serve. When people and their needs are the central focus, how AI should and should not be used becomes clear.
3. **Be accountable:** Using AI doesn't mean absolving humans of responsibility or replacing their efforts entirely. We seek to increase our accountability and environmental stewardship while using AI to augment human capacity, ingenuity and care, instead of automating it away.

Responsible AI Practices

To put these principles into action, we have a set of general practices that cut across our individual and collective use of AI. These practices are the foundation of the more specific guidance below:

1. **Evaluate and Learn at Every Stage:** Whether using AI for big or small tasks, WRI staff should take an evaluative mindset, compare performance to known baselines, prioritize learning, adjust, and share what you learn with colleagues so that the whole organization benefits.
2. **Measure and Manage the Financial and Environmental Cost:** Like other cost and carbon intensive activities such as air travel, commuting, and HVAC, AI tools carry real financial and environmental costs. These should be measured and managed. Individuals should be mindful of their own use, follow the use case guidance below, and participate in periodic usage surveys conducted by WRI's Sustainability Team.
3. **Be Transparent about AI Use:** When AI has shaped your work, say so. Transparency helps colleagues and external audiences calibrate their trust in your product and upholds WRI's credibility and integrity. For more, see the **Disclosure** section below.
4. **Set, Follow, and Give Feedback on Institutional Rules of the Road:** At WRI, we try to provide clear guidance on when and how to use AI responsibly. We ask that staff both follow these rules and give feedback to help improve them over time.
5. **Build Human Capacity and Community:** Use AI as a tool to build your own skills and understanding instead of a shortcut to avoid learning. Share useful approaches and lessons learned with your team and the wider Community.

3. GUIDANCE ON USING AI RESPONSIBLY AT WRI: USE CASES

The decision to use Generative AI should be closely tied to the specific use cases. Generative AI models can work well for one task while performing horribly on others. In addition, the capabilities of these models evolve rapidly over time. Performance on certain tasks can leap ahead, while lagging behind on others.

Currently, you can think of Generative AI as an eager research assistant who is able to analyze large amounts of information and does not mind doing repetitive tasks, but sometimes produces outputs that are not helpful, either because they are inaccurate, generic, or verbose. If you are using Generative AI, your role is to clearly specify the tasks and carefully check the output. **The more complicated the task, the more expertise you need when checking the output. Humans—i.e., you—are ultimately responsible for the work, regardless of how it is produced.**

To guide WRI researchers and communications leads, we offer a framework for rating each use case:

-  **Green Use Cases:** Use Generative AI with appropriate care and **always follow the three rules: (1) keep a human in the loop to review the outputs of the AI; (2) do not share sensitive, confidential, or private information in your prompts to the AI; and (3) do not outsource judgement or taste to AI.**
-  **Yellow Use Cases:** Using Generative AI for this task can be helpful, but only if the user is an expert in the field or topic they are exploring and has the editorial experience to identify and correct gaps and mistakes introduced by AI.
-  **Red Use Cases:** Using Generative AI for this task may cause more harm than good. We currently advise against using Generative AI for red use cases. If you want to experiment with this use case, please see the **Getting Help with Generative AI** section below.

Table 2. Use cases for Generative AI

Use Case Guidance		What You Should Do/Not Do and Why	
	Language/Text		
1	Improving grammar or readability of text	Green	What to do: Use AI to make text simpler and/or more readable while maintaining responsibility for the final content. Be particularly aware of “AI cliches” like excessive use of em dashes, bullets, colon phrases, hyperbolic use of adjectives or adverbs. For more details, see Wikipedia’s Signs of AI Writing page. Why: With care, AI can improve clarity, helping to speed up and improve the drafting process.
2	Brainstorming ideas	Green	What to do: Use AI to generate options and starting points, then test and refine them yourself. Why: AI can suggest useful ideas, but it cannot definitively tell you what’s best for your context.
3	Transcribing interviews	Green	What to do: Use AI to draft transcripts and summaries, then review and correct them before use. As a reminder, do not use AI to transcribe sensitive interviews. Why: Although AI-

			derived transcripts have improved significantly, they can still include errors.
4	Summarizing documents explicitly provided to a model	Yellow	What to do: Only use AI to summarize documents and topics you know well enough to confidently review and correct. Why: AI may miss or underemphasize key points.
5	Translating	Yellow	What to do: Only use AI translations when a subject-matter expert is fluent in the target language and can review and correct the final document. For more details see the WRI Translation Guidelines maintained by the Communications Team. Why: Translation is both art and science, especially around specific technical topics. Subtle differences in language can make a huge difference.
6	Summarizing information generally	Yellow	What to do: Use AI to help summarize information, but verify sources and correct summaries before using or sharing. Why: Accurate and complete summaries are possible but require validation by those who understand the source material.
7	Developing a Literature Review for internal use	Yellow	What to do: Use AI to scan or synthesize literature, ideally with tools built specifically for scientific use cases such as Elicit, Scite, Laser AI or Semantic Scholar AI. Why: Literature reviews require some level of domain knowledge and judgement to get right, but AI can help with aggregation, summarizing, and formatting.
8	Writing short, repetitive, or formulaic text based on a known corpus of information (e.g., generating FAQs from a set of documents, short emails, or a content calendar)	Yellow	What to do: Use AI to write short, formulaic text grounded in source materials, then have an expert review it for accuracy, tone, and clarity. Why: AI can draft efficiently, but the final product can contain inaccuracy or feel generic without appropriate editing.
9	Restructuring text to fit a different format, for example using notes to build a PowerPoint presentation or a policy document to create a training guide	Yellow	What to do: Use AI to reformat notes or text into new structures, then double check for invented details, weak logic, or missing context. Why: Delegating the process of transposition can save significant time but introduces the possibility of errors.
10	Substantive editing, rewriting, or co-drafting of long-form content, for example reports, proposals, or briefing documents	Yellow	What to do: Use AI to help develop draft or restructure and rewrite existing draft while ensuring that human authors remain involved throughout, providing the core ideas, judgment and review at each stage. The final product should reflect human input, not just light editing of AI output. Why: AI can help accelerate writing and improve structure and clarity. But long forming content requires subject matter expertise, judgment and voice that AI cannot reliably

			provide. Without active human involvement, outputs risk being generic or inaccurate.
11	Using Generative AI as a general search engine	Yellow	What to do: Use AI to help find and synthesize information, but check original sources and verify facts. Why: AI-produced search outputs, like any form of search, can be wrong or misleading. Trust but verify.
12	Writing an entire blog, a communications output, or knowledge product	Red	What to do: Do not use AI to draft a full communications or knowledge product piece from scratch. Why: The risk of factual errors, weak reasoning, and noticeable artificial or formulaic language is too high. Doing this undercuts our value as independent experts with good ideas to share.
13	Reviewing research documents sent by others	Red	What to do: If asked to provide your own review and judgment, do not outsource that task to AI without the author's consent. Why: Human judgment and consent matter. If someone is asking you for your review, you should give them your best insight instead of insight from an AI model.
	Coding		
14	Interpreting or documenting existing code	Green	What to do: Use AI to help explain code or draft documentation. Why: Understanding how code work helps software engineers, data scientists, and product managers understand and improve their work. AI is generally good at summarizing available information, code being a solid example.
15	Identifying and fixing bugs in code	Yellow	What to do: Use AI to identify and suggest bugs or fixes only if you're able to test the code and judge the results for yourself. Why: While AI is helpful for troubleshooting, the fixes it implements can add complexity or errors. Human review and care are still critical.
16	Writing code for analysis or small tools	Yellow	What to do: Use tools like Claude Code, Codex, Cursor and/or Replit to generate code for small tools or analysis scripts. Only use AI to generate code that you can write and check yourself. Why: Efficient, effective, and secure code requires expertise, skill, and experience. Taking the time to build your knowledge of engineering and programming languages will make you a better user of AI.
17	Writing test cases	Yellow	What to do: Use AI to generate test cases, edge cases, and failure scenarios as a starting point, but review all assertions carefully before use. Where possible, provide AI with the specification or requirements rather than the existing code, and consider asking it to suggest what to test rather than writing the tests directly. Why: AI-generated tests tend to mirror existing code behavior rather than verify correct behavior, which means bugs can be encoded into tests rather than caught by them. Assertions are often shallow,

			edge cases are frequently missed, and high coverage numbers can create a false sense of security.
18	Building an application or full analysis script	Yellow	What to do: Use AI to generate code for applications or data analysis and/or visualization scripts but audit the output thoroughly. Ensure that the reviewer has the experience needed to identify and fix errors. Why: AI can produce code that works, but still contains bugs, security vulnerabilities or logic errors that may not be immediately obvious, but are still harmful. Without an experienced review, these issues can go undetected.
19	Analyzing qualitative data	Yellow	What to do: Use AI to help organize, transcribe, or summarize qualitative material, but not to define themes or make interpretive judgments without expert review. Why: Qualitative analysis requires contextual knowledge, domain expertise, and methodological judgment that AI lacks. Themes and interpretations generated by AI may miss nuance, reflect bias, or be shaped by patterns in training data rather than the actual meaning of your material.
	Multimedia		
20	Assist with video production tasks including caption generation and script editing	Yellow	What to do: Use AI for tasks like captioning or script polishing, then have an experienced video producer review and correct outputs. For more details, contact WRI's Core Communications team. Why: Video production needs human oversight for quality and style alignment.
21	Creating entire data visualizations or maps	Red	What to do: Do not use image-generation tools for final data visualizations. You may use AI to help write reproducible code that produces visualization, provided a qualified person reviews the data and code (see Use Case 15 above). Why: Final data visualizations should be accurate and reproducible.
22	Creating images or video for internal or external use	Red	What to do: Do not use AI to create images or video. For stock images consult with WRI Core Communications. Why: Generating images and videos is generally more computation and carbon intensive than performing text-based tasks. And images AIs produce may leverage proprietary images lifted from the internet, putting WRI at risk of copyright infringement. Finally, AI generated images create a sense of “unreality” or an “artificial reality” that runs counter to WRI’s commitment to accuracy and authenticity.

GETTING HELP WITH GENERATIVE AI

If you are unsure how to apply these guidelines in your work, want to talk through a new use case not covered here, or would like additional support, we offer the following:

- **Get help Asynchronously:** Join the Data Lab Community’s [AI channel](#) to ask questions or share ideas with colleagues.
- **Demo and Learn with Colleagues Monthly:** The Data Lab organizes the AI Summit on the first Thursday of the month. All are welcome to join and share their work or learn from others. To attend, email evan.tachovsky@wri.org for the invite.
- **Getting Coaching for Your Team:** If you’re interested in an AI coaching session tailored to your team’s needs, email evan.tachovsky@wri.org, laura.malaguzzi@wri.org, and christian.madera@wri.org.

For specific roles, the Data Lab, Research Integrity, and Core Communications have developed collections of more detailed guidance to support responsible use of Generative AI within specific roles:

Roles	Additional Guidance	Maintainer
<i>Communications</i>	• Playbook for Experimenting with AI in Communications for communications professionals.	Core Communications. Contact Christian Madera for details christian.madera@wri.org
<i>Data Science, Product Management, and Software Development</i>	• Generative AI Product Development for product managers, designers, and software developers • Field Guide for AI-Assisted Coding for data scientists or engineers interested in using AI Tools	Data Lab. Contact Evan Tachovsky for details evan.tachovsky@wri.org
<i>Research</i>	• Using AI for literature reviews (Kelly Finke, August 2025) for researchers	Research Integrity. Contact Laura Malaguzzi Valeri for details laura.malaguzzi@wri.org

DISCLOSURE

At WRI, we do not consider Generative AI algorithms to be legitimate co-authors. They are tools. Any use of these tools should be listed in the methods section of WRI’s knowledge products or applications. This information should also be shared with co-authors, reviewers, or editors involved in the production process.

If using a generative AI tool, the authors should include information on:

- The Generative AI platform and/or model type and version used.
- The date it was accessed.
- How the output of the tool was reviewed or adapted, before being included in the final product.

Authors should disclose AI use in:

- Conceptualization of the hypothesis or idea
- Collection, cleaning and analysis of the evidence
- Design or execution of the project
- Development of the project output such as writing a manuscript, building slides, or producing an application

Authors should also disclose how privacy and security were upheld in alignment with [WRI's human subjects protection process](#).

Here is an example of AI disclosure (from Weaver 2024):

Conceptualization: ChatGPT was used to revise research questions; *Data Collection Methods:* ChatGPT was used to create the first draft of the survey instrument; *Data Analysis:* Microsoft Copilot was used to verify identified themes coded from open ended survey responses; *Privacy and Security:* no identifiable data was shared with ChatGPT during the design of this study, only the University of Waterloo institutional instance of Microsoft Copilot was used to analyze any anonymized research data in compliance with University of Waterloo privacy and security policies; *Writing—Review & Editing:* ChatGPT was used in the literature review to provide sentence-level revisions and metaphor options; *Project Administration:* ChatGPT was used to establish a list of tasks and timelines for the study.

4. FAQs

Q: What is the difference between having a human in the loop (green) and having an expert check the work (yellow)?

A: A human should always check AI output. For the yellow category, the human should have specific skills and knowledge relevant to the task at hand, for example knowledge of the scientific domain if you are asking for literature review or knowledge of the programming language if you are asking the model to produce code.

Q: Can I use AI notetaking software like Otter.ai or read.ai to document meetings?

A: AI is useful for summarization, but many third party note takers require deep access to your computer and data, creating security and compliance risk. Staff should use sound judgement and not provide AI notetaking software (or any software) with access to your email, calendar or contacts. Contact WRI's IT team with any question at itsupport@wri.org.

Q: What information is considered sensitive, private or confidential?

A: Many AI platforms use the information you provide as part of their training data, which means that we risk that data becoming part of the public domain. That is problematic for WRI if the information is proprietary or includes personally identifiable information including information shared with us by partners or colleagues in confidence or subject to specific legal agreements. If you have questions about whether the information you are working with is considered sensitive, private, or confidential, ask your manager or reach out to the Legal team.

Q: How can I improve my skills using Generative AI tools?

A: As with any skill, the best way to improve is to practice. Useful resources include Dr. Ethan Mollick’s “[Guide to prompting AI \(for what it is worth\)](#)” blog for basic guidance and Dr. Andrew Ng’s “[ChatGPT Prompt Engineering for Developers](#)” course for more advanced guidance. To chat about Generative AI with other WRI staff, join the [AI channel](#) on Teams.

Q: Are AI detectors reliable?

A: No, not reliably. AI detection tools frequently produce false positives, flagging legitimate human writing as AI-generated, particularly for non-native English speakers. They are also easily defeated by light paraphrasing or editing. Rather than relying on AI detectors, WRI's approach emphasizes transparency, disclosure, and asking the same of others.

Q: What AI tools should I use?

A: There is no single tool that is best for all tasks, and the landscape is evolving rapidly. WRI currently provides staff access to Microsoft Copilot through institutional licenses. Many teams choose to procure other models or tools based on their needs. Claude Code, OpenAI Codex, and Cursor are examples of common tools for software developers. For literature review, domain-specific tools like Elicit or Semantic Scholar AI are preferable to general-purpose chat interfaces. When evaluating any tool, consider whether it:

- It meets WRI's data privacy requirements. Avoid entering sensitive or confidential information into tools without enterprise agreements
- It is appropriate for the use case per Table 2 above
- A WRI-licensed version is available, which typically offers stronger privacy protections than a personal account.

For guidance on specific tools and licenses, contact the IT and Data Lab teams.

Have a question that is not answered in this document? Please ask it [here](#).

5. REFERENCES

“ChatGPT Cheat Sheet: Complete Guide for 2023.” 2023. *TechRepublic* (blog). July 21, 2023.

<https://www.techrepublic.com/article/chatgpt-cheat-sheet/>.

Acemoglu, D., Dingwen K., Ozdaglar, A. 2026. AI, Human Cognition and Knowledge Collapse. Working Paper No. 34910. Working Paper Series. National Bureau of Economic Research, February.

<https://doi.org/10.3386/w34910>.

Azaria, A, Mitchell, T. 2023. “The Internal State of an LLM Knows When It’s Lying.” *Findings of the Association for Computational Linguistics: EMNLP 2023*, 967–76.

<https://doi.org/10.18653/v1/2023.findings-emnlp.68>.

Brynjolfsson, E., Li, D., Raymond, L. 2025. Generative AI at Work, *The Quarterly Journal of Economics*,

<https://doi.org/10.1093/qje/qjae044>

Budzyń, K., M. Romańczyk, D. Kitala, P. Kołodziej, M. Bugajski, H. O. Adami, J. Blom, M. Buszkiewicz, N. Halvorsen, C. Hassan, T. Romańczyk, Ø. Holme, K. Jarus, S. Fielding, M. Kunar, M. Pellise, N. Pilonis, M. F. Kamiński, M. Kalager, M. Bretthauer, and Y. Mori. 2025. “Endoscopist Deskilling Risk after Exposure to

Artificial Intelligence in Colonoscopy: A Multicentre, Observational Study.” *The Lancet Gastroenterology & Hepatology* 10 (10): 896–903. [https://doi.org/10.1016/S2468-1253\(25\)00133-5](https://doi.org/10.1016/S2468-1253(25)00133-5)

Chen, S. (2025, March 5). How much energy will AI really consume? The good, the bad and the unknown. *Nature*. <https://www.nature.com/articles/d41586-025-00616-z>

Cui, Z., Demirer, M., Jaffe, S., Musolff, L., Peng, S., & Salz, T. 2026. The Effects of Generative AI on High-Skilled Work: Evidence from Three Field Experiments with Software Developers. *Management Science*, ahead of print, February 27. <https://doi.org/10.1287/mnsc.2025.00535>.

Dell’Acqua, F., E. McFowland, E. Mollick, et al. 2026. “Navigating the Jagged Technological Frontier: Field Experimental Evidence of the Effects of Artificial Intelligence on Knowledge Worker Productivity and Quality.” *Organization Science*, March 11. doi:[10.1287/orsc.2025.21838](https://doi.org/10.1287/orsc.2025.21838).

Dis, E.A.M. van, J. Bollen, W. Zuidema, R. van Rooij, and C.L. Bockting. 2023. “ChatGPT: Five Priorities for Research.” *Nature* 614 (7947): 224–26. doi:[10.1038/d41586-023-00288-7](https://doi.org/10.1038/d41586-023-00288-7).

Dodge, J., Gururangan, S., Card, D., Schwartz, R., & Smith, N. A. 2023. Measuring the Carbon Intensity of AI in Cloud Instances. *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.

Dwivedi, Y.K., N. Kshetri, L. Hughes, E.L. Slade, A. Jeyaraj, A.K. Kar, A.M. Baabdullah, et al. 2023. “Opinion Paper: ‘So What If ChatGPT Wrote It?’ Multidisciplinary Perspectives on Opportunities, Challenges and Implications of Generative Conversational AI for Research, Practice and Policy.” *International Journal of Information Management* 71 (August): 102642. doi:[10.1016/j.ijinfomgt.2023.102642](https://doi.org/10.1016/j.ijinfomgt.2023.102642).

Elsworth, C., Keguo H., Patterson, D. et al. 2025. Measuring the Environmental Impact of Delivering AI at Google Scale. arXiv:2508.15734. Preprint, arXiv, August 21. <https://doi.org/10.48550/arXiv.2508.15734>.

Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y., Chen, D., Dai, W., Chan, H. S., Madotto, A., & Fung, P. 2023. Survey of Hallucination in Natural Language Generation. *arXiv preprint arXiv:2202.03629*. <https://doi.org/10.48550/arXiv.2202.03629>

Malmqvist, L. 2025. Sycophancy in Large Language Models: Causes and Mitigations. In: Arai, K. (eds) *Intelligent Computing. CompCom. 2025. Lecture Notes in Networks and Systems*, vol 1426. Springer, Cham. https://doi.org/10.1007/978-3-031-92611-2_5 Noy, S., and W. Zhang. 2023. “Experimental Evidence on the Productivity Effects of Generative Artificial Intelligence.” *Science* 381 (6654): 187–92. doi:[10.1126/science.adh2586](https://doi.org/10.1126/science.adh2586).

National Institute of Standards and Technology (NIST). 2024. *Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile*. NIST AI 600-1. Gaithersburg, MD: U.S. Department of Commerce. <https://doi.org/10.6028/NIST.AI.600-1>

Saenko, K. 2023. “Is Generative AI Bad for the Environment? A Computer Scientist Explains the Carbon Footprint of ChatGPT and Its Cousins.” *The Conversation* (blog). May 23, 2023. <http://theconversation.com/is-generative-ai-bad-for-the-environment-a-computer-scientist-explains-carbon-footprint-of-chatgpt-and-its-cousins-204096>.

Shen, J. H., and A. Tamkin. 2026. “How AI Impacts Skill Formation.” *arXiv preprint arXiv:2601.20245*. <https://doi.org/10.48550/arXiv.2601.20245>

Shumailov, I., Z. Shumaylov, Y. Zhao, N. Papernot, R. Anderson, and Y. Gal. 2024. "AI Models Collapse When Trained on Recursively Generated Data." *Nature* 631: 755–759. <https://doi.org/10.1038/s41586-024-07566-y>.

Stokel-Walker, C., and R. Van Noorden. 2023. "What ChatGPT and Generative AI Mean for Science." *Nature* 614 (7947): 214–16. doi:[10.1038/d41586-023-00340-6](https://doi.org/10.1038/d41586-023-00340-6).

Wang, Z.Z., Shao, Y., Shaikh, O., Fried, D., Neubig, G., and Yang, D. 2025. "How Do AI Agents Do Human Work? Comparing AI and Human Workflows Across Diverse Occupations." arXiv preprint arXiv:2510.22780. <https://doi.org/10.48550/arXiv.2510.22780>

Weaver, K. D. 2024. *The artificial intelligence disclosure (AID) framework: An introduction*. *C&RL News*, 85(10), 407-411. <https://crln.acrl.org/index.php/crlnews/article/view/26548>

Weisz, J.D., M. Muller, J. He, and S. Houde. 2023. "Toward General Design Principles for Generative AI Applications." In *Joint Proceedings of the IUI 2023 Workshops: HAI-GEN, ITAH, MILC, SHAI, SketchRec, SOCIALIZE*, edited by A. Smith-Renner and P. Taele, 3359:130–44. CEUR Workshop Proceedings. Sydney, Australia: CEUR. <https://ceur-ws.org/Vol-3359/#paper14>.

Whittlestone, J., & Clarke, S. 2023. AI Challenges for Society and Ethics. In *The Oxford Handbook of AI Governance*. Oxford University Press

Zou, A., Wang, Z., Carlini, N., Nasr, M., Kolter, J. Z., & Fredrikson, M. 2023. Universal and Transferable Adversarial Attacks on Aligned Language Models. *arXiv preprint arXiv:2307.15043*. <https://doi.org/10.48550/arXiv.2307.15043>

6. POLICY UPDATE RECORD

Date	Summary of Changes Made	Approved By
2026-04-28	Significant changes include: - Added a third Green Use Case rule. The two existing rules for Green use cases (keep a human in the loop; do not share sensitive information) have been supplemented with a third: "do not outsource judgement or taste to AI." - Overhauled the Principles section. The OECD AI Principles have been replaced with WRI's own Responsible AI Principles and Practices, organized around three institutional principles (Be Curious, Be User Centered, Be Accountable) and five individual practices (Evaluate and Learn; Measure and Manage Cost; Be Transparent; Follow and Give Feedback on Institutional Rules; Build Human Capacity and Community). - Expanded the "How Does Generative AI Work" section to cover reasoning models and agentic modes, and updated platform/model footnote to reflect the April 2026 landscape (GPT-5, Claude-4, Llama-4, Mistral). - Updated the evidence base on AI efficacy. Added new findings on sycophancy (Malmqvist 2025), deskilling risks	Data Lab Global Director, Research Integrity Global Director with feedback from the GET AI Working Group, the Applied AI Group, and staff via the AI Summit.

	(Acemoglu et al. 2026), and automation-verification tradeoffs (Wang et al. 2025). Removed the "quality degradation over time" finding from the emerging consensus section. - Restructured the use case table, adding a "What You Should Do/Not Do and Why" column in place of the brief guidance notes. Specific use case changes include: added "Reviewing documents sent by others" as a new Red use case; changed "Build an entire application or analysis script" from Red to Yellow with conditions; added "Analyzing qualitative data" as a new Yellow use case; changed "Creating data visualizations" from Yellow to Red for final outputs (while permitting AI-assisted code to produce reproducible visualizations); and expanded the images use case to include video. - Added links to role-specific guidance for Communications; Data Science, Engineering, Product, and Research. - Significantly expanded the Disclosure section, adding disclosure categories (conceptualization, data collection/analysis, design/execution, output development, privacy/security) and a worked example based on Weaver (2024). - Replaced the "Experimenting with New Use Cases" section with a new "Getting Help with Generative AI" section providing structured channels for async support, monthly learning sessions, and team coaching. - Added two new FAQs: "Are AI detectors reliable?" and "What AI tools should I use?" - Minor edits to improve readability and coherence throughout, including updated references and hyperlinks. AI Disclosure: <i>Writing, Review and Editing:</i> Claude (Anthropic, claude-sonnet-4-6, accessed March and April 2026) was used to suggest edits to improve the clarity and readability of text across the document. <i>Writing, Development:</i> Claude was used to transpose and reformat existing content from one structural format to another, including restructuring the use case table and transposing WRI's principles and practices from previously published documents. <i>Analysis:</i> Claude was used to review the overall guidance document and provide an assessment of its strengths, weaknesses, and internal coherence. <i>Project Administration:</i> Claude was used to compare versions and draft the change log entry above	
--	--	--

	All AI-generated suggestions were manually reviewed and edited by the authors, and all work presented here is their responsibility alone.	
2025-04-08	Significant revision to guidelines. Specific changes include: - Updates to the “How Does Generative AI Work” section to ensure it accurately describes current approaches. - Updates to the “What do we Know About Generative AI Efficacy” section to restructure it for clarity and update the literature review with new research. - Updates to the “Common Risk and Opportunities” to rank the risk/opportunities in line with WRI’s work. - Adjusting risk ratings based on current capabilities. The biggest shift was moving “Using Generative AI as a General Search Engine” from Red to Yellow since nearly all models now have ways to access real time data and/or are directly integrated with commercial search engines. - Clarification of vocabulary differentiating different aspects of Generative AI. We use “tools” as a generic term, “platforms” to refer to systems like ChatGPT with a user interface and “models” to refer to the underlying algorithms and data. - Edits to text to improve readability across the document Note that Generative AI (ChatGPT 4.5 access March-April 2025) was used to (1) analyze the text and suggest edits to improve clarity, (2) suggest edits to the description of Generative AI technology, and (3) suggest additions to the literature review. All suggested edits and additions were manually reviewed by the authors and all work presented here is their responsibility alone.	Global Directors for Data Lab, Research Integrity and Digital Communications
2023-09-18	Original draft published	